

Assignment - 6

Modifying Pintos to run user programs with arguments and system calls.

---- GROUP 7 ----

>> Fill in the names, roll numbers and email addresses of your group members.

Gurjot Singh Suri 17CS10058 anshusuri123@gmail.com

Kumar Abhishek 17CS10022 ohyesabhi398@gmail.com

---- PRELIMINARIES ----

>> If you have any preliminary comments on your submission, notes for the
>> TAs, or extra credit, please give them here.

>> Please cite any offline or online sources you consulted while
>> preparing your submission, other than the Pintos documentation, course
>> text, lecture notes, and course staff.

ARGUMENT PASSING =====

---- DATA STRUCTURES ----

>> A1: Copy here the declaration of each new or changed `struct' or
>> `struct' member, global or static variable, `typedef', or
>> enumeration. Identify the purpose of each in 25 words or less.

```
static bool setup_stack (void **esp, char *file_name);
```

We didn't declare new struct or change struct member or use global or static variable for argument passing. We changed static bool setup_stack (void **esp) to static bool setup_stack (void **esp, char *file_name), the new argument is used for parsing the arguments using strtok_r() which are used to fill stack.

---- ALGORITHMS ----

>> A2: Briefly describe how you implemented argument parsing. How do
>> you arrange for the elements of argv[] to be in the right order?
>> How do you avoid overflowing the stack page?

in load function :

stack is set up and arguments are filled up by calling `setup_stack()` with stack pointer and filename's copy as arguments.

in `setup_stack()` :

`esp` (stack pointer) is set to point to the `PHYS_BASE` (begin of stack)

check to see if `install_page` fails to return

a copy of `file_name` from the command line is split on spaces using `strtok_r()` and `argc` is set appropriately until token is `NULL`, then `argv[]` is created which is of size `argc`

`argv` is filled up with arguments by splitting the copy of filename again

`arg_ptr[]` array is created to store addresses of arguments of the stack, its size is `argc`

arguments are stored on the stack in reverse order using `argv[]`

for each argument, the stack pointer is decremented, then stored on the stack, the stack address is stored in `arg_ptr[]`, `total_leng` is incremented by `arg` size

word alignment is added and then `esp` (stack pointer) is updated by checking whether `total_leng` is a multiple of 4
null character is added and `esp` is updated by subtracting the `sizeof(char*)`

argument addresses are added to the stack in reverse order like before from `arg_ptr[]` and then `esp` is updated by subtracting the size of `char *`

this is repeated for all the `argc` addresses in `arg_ptr[]`

size of `char**` is subtracted from `esp` (stack pointer)

stack address of the first argument is added on the stack, i.e. the last address stored on stack.

fake return address is added and `esp` is updated

How do you arrange for the elements of `argv[]` to be in the right order?

The command is parsed to get the number of arguments in `argc`,

then we construct an `argv` array with datatype `char *`, and of size `argc`,

then arguments are stored in `argv` by again parsing the same command,

then arguments are stored on the stack in the right to left order, addresses in `arg_ptr[]`

by traversing the `argv[]` array in reverse order. Wordalign, the stack addresses from `arg_ptr[]` on the stack, address of first argument on stack, `argc`, null char and the fake return address are stored.

How do you avoid overflowing the stack page?

Checking the `esp` pointer is not required until it fails, like add another `argv` element, when necessary.

Overflowing is dealt by letting it fail and then it is handled by the page fault exception,

which then further calls `sys_exit(-1)` for the running thread whenever the address is invalid.

Also the process is terminated if it provides too much arguments.

---- RATIONALE ----

>> A3: Why does Pintos implement `strtok_r()` but not `strtok()`?

strtok_r() and strtok() has only 1 difference which is that the save_ptr (placeholder) in strtok_r() is provided by the caller. Commands are separated into command line (executable name) and arguments by kernel of Pintos. Addresses of the arguments are to be stored somewhere such that they can be reached later such that we are sure if there were more than one thread call strtok()_r each thread have their own pointer (save_ptr) which is independent from the caller and can remember their position.

>> A4: In Pintos, the kernel separates commands into a executable name
>> and arguments. In Unix-like systems, the shell does this
>> separation. Identify at least two advantages of the Unix approach.

It is cleaner to separate the executable name from the arguments before passing it off to the kernel, since they represent different things. It should be parsed by a user program rather than by kernel.

Validation of the input is done by the shell more safely than by the kernel. If someone entered a very large amount of text, perhaps it would cause the kernel a problem if the kernel tried to parse it, whereas if the shell takes care of it, worst case is the shell crashes.

It can separate the commands, it can do advanced pre-processing, acting more like an interpreter not only an interface. Like passing more than 1 set of command line at a time

SYSTEM CALLS =====

---- DATA STRUCTURES ----

>> B1: Copy here the declaration of each new or changed 'struct' or
>> 'struct' member, global or static variable, 'typedef', or
>> enumeration. Identify the purpose of each in 25 words or less.

In syscall.h:

```
/*we add new struct for storing open file descriptors of thread in fd_list*/
struct fd_elem
{
    int fd;                /*file descriptors ID*/
    struct file *myfile;    /* the real file*/
    struct list_elem elem;  /*list elem to add fd_element in fd_list*/
}
/*Global variable Used to ensure that only one process at a time
```

is executing file system code. *

```
struct lock file_lock;
```

In thread.h

```
/*we add new struct for storing children of thread in child_list*/
struct child_elem
{
    bool waited;          /*to check if wait() is called before?*/
    struct list_elem elem; /*list elem used to add in child_list */
    bool load_success;     /*to check if load success*/
    int exit_status;       /*the status the child thread exit with*/
    int cur_status;        /*the child thread current status*/
    int child_pid;         /*pid of this child*/
    struct semaphore sem;  /*parent wait child to load or exit*/
};

/*adding some change in struct thread*/
struct thread{
    struct list child_list; /*list of children this thread have*/
    struct thread * parent; /*pointer to this thread's dad*/
    /*file file descriptors*/
    struct list fd_list;    /*list of file descriptors*/
    int fd_size;            /*size of the file descriptors*/
    /*executable file should not edited while running*/
    struct file *exec_name; /*executable held by this thread*/
};
```

>> B2: Describe how file descriptors are associated with open files.

>> Are file descriptors unique within the entire OS or just within a

>> single process?

Within a single process file descriptors are unique. File_descriptors list fd_list (list of struct fd_elem) is being tracked by process, as well as its next available descriptor number in fd_size which is incremented every time. Our fd_elem struct is what associates the file descriptor numbers with the corresponding file

---- ALGORITHMS ----

>> B3: Describe your code for reading and writing user data from the

>> kernel.

Addresses while reading or writing are validated by calling check_valid_ptr to check if it is not null and that it points below PHYS_BASE using is_user_vaddr

and if it is mapped using `pagedir_get_page` functions, otherwise it calls `sys_exit(-1)`;
Addresses having syscall number are validated and obtained. Then the three arguments are obtained by dereferencing the addresses after validating them. Addresses `buffer` and `buffer+size-1` are also validated by calling `check_valid_ptr`. If a page fault is caused by an invalid user pointer, it is handled in `userprog/exception.c` and then terminates the process with `sys_exit(-1)`.
Read is done by calling read function, write is done by calling write function, if the addresses are valid,

in read :

`int read(int fd, void *buffer, unsigned size)` is called
`lock_acquire(&file_lock)` is called to ensure at a certain time only one process is using the file.
if `fd` is greater than 0 we call `extract_fd()` which iterates on the `fd_list` of the current thread and get the file which have the same `fd`
if not found return `NULL`
`file_read()` which is present in `file.c` is called and
its return value - the number of bytes actually read or returns -1 in case of an error, i.e set `eax` to return val.
`lock_release(&file_lock)` is called before any return.

in write :

`int write (int fd, const void *buffer, unsigned size)` is called
`lock_acquire(&file_lock)` is called to ensure at a certain time only one process is using the file.
`fd` is checked if it is equal to 1 then `putbuff()` is called to write to console and return the size else `extract_fd()` is called which iterates on the `fd_list` of the current thread and used to get the file which have the same `fd` if not found `NULL` is returned.
`file_write()` is called which is present in `file.c` and it returns the value of the number of bytes actually written or -1 in case of error, i.e set `eax` to return val.
`lock_release(&file_lock)` is called before any return.

>> B4: Suppose a system call causes a full page (4,096 bytes) of data
>> to be copied from user space into the kernel. What is the least
>> and the greatest possible number of inspections of the page table
>> (e.g. calls to `pagedir_get_page()`) that might result? What about
>> for a system call that only copies 2 bytes of data? Is there room
>> for improvement in these numbers, and how much?

If the user space data spans over two pages then it would take max 2 calls to `pagedir_get_page` because it may be held over two pages. If it's all on one page then it would take 1 call to `pagedir_get_page`.

If a system call causes a full page of data to be copied, the least possible number of inspections of the page table is 1 and the greatest possible number

is 2, this depends on if the data spans 1 page or 2 page.

For a system call that copies 2 bytes of data, the least possible number of inspections of the page table is 1 and the greatest possible number is 2 too (although much more likely to only require 1).

There's no way to avoid the second lookup so this can not be improved.

>> B5: Any access to user program memory at a user-specified address
>> can fail due to a bad pointer value. Such accesses must cause the
>> process to be terminated. System calls are fraught with such
>> accesses, e.g. a "write" system call requires reading the system
>> call number from the user stack, then each of the call's three
>> arguments, then an arbitrary amount of user memory, and any of
>> these can fail at any point. This poses a design and
>> error-handling problem: how do you best avoid obscuring the primary
>> function of code in a morass of error-handling? Furthermore, when
>> an error is detected, how do you ensure that all temporarily
>> allocated resources (locks, buffers, etc.) are freed? In a few
>> paragraphs, describe the strategy or strategies you adopted for
>> managing these issues. Give an example.

We first check the frame's stack pointer esp to be valid using `check_valid_ptr` then dereference it if valid and check the system call number.

Secondly, we check the pointers to arguments using `check_valid_ptr()` and obtain them, if the arguments are themselves pointers like `char *`, we also validate them, (like in `OPEN` and `CREATE`).

In case of `WRITE` and `READ` system calls, we additionally also validate `buffer+size-1` address by calling `check_valid_ptr()` to check the span.

If any pointer is invalid we call `sys_exit(-1)` that call `thread_exit` which in turn call `process_exit` where we release all the resources acquired by the thread.

If an invalid user pointer still causes a "page fault", it is handled in `userprog/exception.c` and terminates the process with `sys_exit(-1)`.

We ensure all resources are freed by calling `sys_exit(-1)` that call `thread_exit` which in turn call `process_exit` where we release all the resources acquired by the thread.

Example: Dealing with bad-ptr during `READ` (for example), so we call

check_valid_ptr(), If the bad_ptr is NULL or greater than PHYS_BASE this function catch it and call sys_exit(-1)
Then, we call check_valid_ptr() to check if the buffer (second argument) spans in user page by checking buffer and buffer+size using check_valid_ptr()

---- RATIONALE ----

>> B6: Why did you choose to implement access to user memory from the
>> kernel in the way that you did?

We chose to implement access to user_memory by verifying the validity of a user-provided pointer by calling check_valid_ptr which checks if it is not null and that it points below PHYS_BASE using is_user_vaddr and if it is mapped using pagedir_get_page functions, otherwise it calls exit(-1). Only then we dereference it.

We chose it because it is a simpler way to handle user memory access. It's more difficult to handle freeing of resources if an invalid pointer causes a page fault, because there's no way to return an error code from a memory access. We ensure all resources are freed by calling sys_exit(-1) that call thread_exit which in turn call process_exit where we release all the resources acquired by the thread. It decreases some performance in case of valid pointers but handling memory access is easier.

>> B7: What advantages or disadvantages can you see to your design
>> for file descriptors?

Advantages:

Regardless of whether our file descriptors are created by pipe or open, the same structure can store the necessary information, and be used in essentially the same way.

Dynamic list is used to store the current process open files, so we need not worry about its size. Kernel is aware of all the open files, which gains more flexibility to manipulate the opened files. Thread-struct's space is minimized.

Disadvantages:

Accessing a file descriptor is $O(n)$, where n is the number of file descriptors for the current thread (have to iterate through the entire fd list). Could be $O(1)$ if they were stored in an array, but then array size would not be fixed.

The inherits of open files opened by a parent require extra effort to be implement. Consumes kernel space, user program may open lots of files to crash the kernel.